Large-scale Clustering of People Diagnosed with Parkinson's Disease using Acoustic Analysis of Sustained Vowels: Findings in the Parkinson's Voice Initiative Study

Athanasios Tsanas[1] [a] and Siddharth Arora[2] [b]

[1]*Usher Institute, Edinburgh Medical School, University of Edinburgh, U.K.*
[2]*Department of Mathematics, University of Oxford, U.K.*

Keywords: Acoustic Analysis, Clustering, Parkinson's Disease, Parkinson's Voice Initiative (PVI).

Abstract: The heterogeneity of symptoms in Parkinson's Disease (PD) has motivated investigating PD subtypes using cluster analysis techniques. Previous studies investigating PD clustering have typically focused on symptoms assessed using standardized clinical evaluations and patient reported outcome measures. Here, we explore PD subtype delineation using speech signals. We used data from the recently concluded Parkinson's Voice Initiative (PVI) study where sustained vowels were solicited and collected under non-controlled acoustic conditions. We acoustically characterized 2097 sustained vowel /a/ recordings from 1138 PD participants using 307 dysphonia measures which had previously been successfully used in applications including differentiating healthy controls from PD participants, and matching speech dysphonia to the standard PD clinical metric quantifying symptom severity. We applied unsupervised feature selection to obtain a concise subset of the originally computed dysphonia measures and explored hierarchical clustering combined with 2D-data projections using t-distributed stochastic neighbor embedding to facilitate visual exploration of PD subgroups. We computed four main clusters which provide tentative insights into different dominating speech-associated pathologies. Collectively, these findings provide new insights into the nature of PD towards exploring speech-PD data-driven subtyping.

1 INTRODUCTION

Parkinson's Disease (PD) is a progressive neurodegenerative disorder with continuously increasing prevalence rates and growing burden for national health systems (Dorsey et al., 2013). In 2016 there were approximately 6.1 million people reportedly diagnosed with PD compared to 2.5 million people in 1990 (GBD, 2018). The primary PD symptom constellation comprises tremor, rigidity, bradykinesia, and postural stability. These fit within the broader spectrum of variable factors including motor, cognitive, and neuropsychiatric symptoms (Olanow, Stern, Sethi 2009). PD is well reported as a largely heterogeneous disease, which is further accentuated with considerable heterogeneity in individual patient symptom severity trajectories (Fereshtehnejad et al., 2015).

Assigning PD participants into subtypes is clinically important since homogeneous groups exhibit stronger clinical symptom manifestation and potentially stronger genetic coherence. Therefore, understanding different PD subtypes may lead to new insights towards involved biological pathways, which in turn may lead to better-informed, targeted treatment strategies. In practice, PD group membership may be achieved using some predefined clinical intuition and criteria such as age onset and dominating symptoms. Data-driven approaches to delineate PD subtypes have received increasing attention in the research community over the last few years (Lewis et al., 2005; Selikhova et al., 2009; Lawton, 2018). Indicative examples include using clinico-pathological characteristics (Selikhova et al., 2009), standardized clinical instruments to assess motor, non-motor, and cognitive domains (Lawton, 2018), or sensor-based gait pattern analysis (Nguyen

[a] https://orcid.org/0000-0002-0994-8100

[b] https://orcid.org/0000-0001-6499-6941

et al. 2019). The use of different modalities or clinical instruments to assess symptoms can potentially provide new insights, but makes comparisons across studies particularly challenging and may explain discrepancies in the reported PD subtypes.

Crucially for the purposes of this study, speech is very strongly associated with overall PD symptom severity as assessed using standardized clinical metrics (Tsanas, 2019) and 29% of people diagnosed with PD consider it one of their most debilitating symptoms (Hartelius and Svensson, 1994). Recent studies have demonstrated the potential of speech signals and in particular sustained vowel /a/ phonations in PD applications, e.g. to (1) differentiate Healthy Controls (HC) from people diagnosed with PD with almost 99% accuracy (Tsanas et al., 2012), (2) accurately replicate the Unified Parkinson's Disease Rating Scale (UPDRS) (Tsanas et al., 2011), which is the standard clinical tool to provide an overall PD symptom assessment, and (3) automatically assess voice rehabilitation (Tsanas et al., 2014a). Recent work has also demonstrated the potential of speech signals towards distinguishing people with Leucine-Rich Repeat Kinase 2 (LRRK2) associated PD, idiopathic PD, and HC (Arora et al., 2018). Moreover, speech articulation kinematic models to characterize PD dysarthria and provide insights into the underlying vocal production mechanism have been developed (Gomez et al., 2019). Collectively, these studies and many others demonstrate the enormous potential of using speech signals in the context of PD analysis.

The aforementioned diverse problems rely on the existence of clinical labels and belong to the *supervised learning* paradigm. In situations where clinical labels (i.e. outcomes of interest) are not available, researchers typically resort to *unsupervised learning* methods for data exploration. These exploratory methods aim to decipher hidden patterns in the data or provide the means towards understanding the internal data structure e.g. with *cluster analysis* methods (Hastie, Tibshirani, Friedman, 2009). Cluster analysis aims to group together "similar" data samples (also known as *objects* in the statistics parlance) and in distinct groups data samples which are "different". There are numerous strategies and algorithms for cluster analysis where the central notion is the concept of computing similarity amongst objects (see Hastie, Tibshirani, Friedman, 2009; Duda, Hart, Stork, 2001). In simple terms, each object is assigned (or probabilistically assigned) cluster membership. The resulting outputs of cluster analysis are known as clusters or groups, and are often referred to as derived subtypes in clinical applications.

Most studies aiming to report PD subtypes rely on standard cluster analysis methods and in particular k-means (e.g. Lewis et al., 2005; Lawton et al., 2018), which is one of the simplest approaches but which is known to have some fundamental drawbacks (Hastie, Tibshirani, Friedman, 2009; Duda, Hart, Stork, 2001). Additional considerations in cluster analysis include how to select a robust feature subset in an unsupervised feature selection framework (Dy and Brodley, 2004), potentially standardizing variables or introducting weights for different variables, and finally validating findings. Unfortunately many of the finer details in the application of the end-to-end cluster analysis methodology in clinical studies are frequently not reported. For an overview of this field (albeit using a different clinical application as an exemplar), including highlighting shortfalls and suggestions for best practice when reporting clustering results we refer to Horne et al. (2020).

The aim of this study is to explore speech-PD data-driven subtyping using cluster analysis methods and provide tentative new insights into the nature of PD speech symptoms. Towards this aim we acoustically characterize sustained vowel /a/ phonations, determine a subset of dysphonia measures using unsupervised feature selecton, and experiment with different cluster analysis and data visualization tools.

2 DATA

The PVI study solicited phone calls from participants across seven major geographical locations (Argentina, Brazil, Canada, Mexico, Spain, USA, and the UK). People were requested to call a dedicated phone number and contribute (1) basic demographic information (age, gender), (2) self-report whether they had been clinically diagnosed with PD, and (3) two sustained vowel /a/ phonations. Following standard voice assessment protocols participants were instructed to sustain vowel /a/ for as long and as steadily as possible (Titze, 2000). Recordings were sampled at 8 kHz and stored on secure servers hosted by Aculab.

In this study we only processed data from the PD participants to investigate PD subtypes and discarded data contributed by HC. Furthermore, we focus only on the data from the US cohort (geographic location with most data) to simplify analysis and avoid language confounds which might be otherwise reflected in the clustering results. In total, we processed 2097 sustained vowel /a/ phonations from 1138 PD participants (605 males) with age (mean ± standard deviation): 63.7±10.8 years. For further details on the PVI study

we refer to our previous work (Arora, Baghai-Ravary, Tsanas, 2019; Tsanas and Arora, 2019).

3 METHODS

3.1 Data Pre-processing

We developed a speech recognition software which automatically transcribed the participants' responses over the phone regarding age, gender, and self-reported PD assessment. When the automated speech recognition algorithm had less than 90% confidence regarding the participants' responses, the recordings were aurally inspected. Furthermore, we developed signal processing tools to screen out non-usable recordings e.g. with excessive background noise.

For further details please see (Arora, Baghai-Ravary, Tsanas, 2019).

3.2 Acoustic Characterization of Sustained Vowel /a/ Phonations

We used the Voice Analysis Toolbox (freely available on the first author's website: https://www.darth-group.com/software) to acoustically characterize each sustained vowel /a/ phonation using 307 dysphonia measures. The toolbox includes a range of widely used dysphonia measures which have been developed specifically to characterize sustained vowel /a/ phonations, and has been extensively validated in PD applications (Tsanas et al., 2010; Tsanas et al., 2011; Tsanas et al., 2012; Tsanas, 2012; Tsanas et al., 2014a; Arora, Baghai-Ravary, Tsanas, 2019), and other voice-related applications (Tsanas and Gomez-Vilda, 2013; San Segundo, Tsanas, Gomez-Vilda, 2017). For the underlying rationale, conceptual basis and physiological background, as well as the algorithmic expressions for the computation of the dysphonia measures we refer to (Tsanas, 2012; Tsanas, 2013). A key component in speech signal analysis which is frequently a prerequisite for the computation of more advanced dysphonia measures is the fundamental frequency (F0), and in particular its time-varying property also known as F0 *contour*. We used the SWIPE algorithm (Camacho and Harris, 2008), which we had previously demonstrated is the most accurate F0 estimation algorithm in sustained vowel /a/ phonations (Tsanas et al., 2014b). Overall, applying the speech signal processing algorithms to each of the phonations in the study resulted in a 2097×307 feature matrix which was subsequently mined to determine possible cluster solutions. All features are continuous random variables.

Before using the 307 features in the subsequent stages we linearly scaled each feature to be in the range [0, 1] so that no feature dominates others, in accordance to the standard rule of thumb for distance-based machine learning algorithms (Bishop, 2006).

3.3 Unsupervised Feature Selection

A high dimensional dataset may increase the noise to signal ratio and obscure data structure and pattern recognition algorithms. This standard problem is known as the *curse of dimensionality* and is often detrimental for the performance of machine learning algorithms (Guyon et al. 2006; Hastie, Tibshirani, Friedman, 2009). According to the general *principle of parsimony*, it is desirable to develop a predictive model which at the same time is as simple as possible, i.e. via reducing the dimensionality of the input space. This approach is known as *dimensionality reduction*, and can be achieved either by *feature transformation* (transforming the features to populate a new, lower dimensional space), or by *feature selection* (choosing a subset of features from the original feature set). The latter is typically preferred in clinical settings because it is desirable to retain the interpretability of the original features (Guyon et al., 2006; Tsanas, Little, McSharry, 2013).

In supervised learning frameworks, feature selection can be wrapped around a well-defined objective function capitalizing on the provided labels. Feature selection in unsupervised learning setups is less well defined and therefore more challenging (Dy and Brodley, 2004). The aim is identifying informative features supporting complex structures embedded in the high-dimensional space, as Dy and Brodley (2004) suggest: "*The goal of feature selection for unsupervised learning is to find the smallest feature subset that best uncovers "interesting natural" groupings (clusters) from data according to the chosen criterion.*"

Here, we used the algorithmic approach endorsed by Yao et al. (2015) called *i-Detect* to select informative features where the identified feature subspace has the following property: the difference between the total volume of the space spanned by the selected feature subset and the sum of the volumes of clusters in the embedded manifolds is maximized. The i-*Detect* algorithm has two free hyper-parameters which need to be optimized: the kernel width, and the regularization parameter. Given that the algorithm is not sensitive to the choice of the kernel width (Yao et al. 2015), we focused only on experimenting with the selection of the regularization parameter.

Ultimately, the output of this unsupervised feature selection algorithm is a feature weight vector where many of the features are assigned to zero weighting and hence can be eliminated. The computed weights are then used to rank the original features and decide on an appropriate cut-off.

3.4 Clustering

Clustering falls under the unsupervised learning category and aims to provide some insight into the structure of the data to and group objects based on the similarity of the provided features. The output of a clustering algorithm indicates the (probabilistic) cluster membership of each object into the possible clusters. There are many clustering algorithms in the research literature, each with shortcomings and different strategies to optimize performance.

In this study, we used hierarchical clustering which is a popular cluster analysis method that has often been successfully used in diverse applications (Hastie, Tibshirani, Friedman, 2009). Unlike other competing cluster analysis methods such as k-means, hierarchical clustering does not require pre-specifying the number of clusters in the data. Hierarchical clustering constructs a dendrogram to represent the data in a tree-based form, which intuitively depicts how objects are grouped in the form of different levels. The tree is recursively split to form new clusters, aiming to maximize the between group dissimilarity. For further background details we refer readers to (Duda, Hart, and Stork 2004).

We used hierarchical clustering with Ward's linkage to cluster both the original high-dimensional data and the lower-dimensional representation obtained following unsupervised feature selection. The number of clusters was determined following visual inspection of the dendrogam in accordance with Sheaves et al. (2016). In essence, we aim to find a cut-off where there is considerable dissimilarity difference between successive levels.

3.5 Data Visualization

We applied the t-distributed Stochastic Neighbor Embedding (t-SNE) algorithm (van der Maaten and Hinton, 2008) to visualize the data structure embedded in the high-dimensional space (using the original 307-dimensional space and also the feature space spanned with the selected features). The resulting 2D data representation can potentially provide new insights following visual inspection and can also be used to visually assess the cluster analysis results.

4 RESULTS

Figure 1 presents the dendrogram when using the original high-dimensional feature set prior to feature selection. Based on visual inspection, we decided to opt for six clusters (highlighted with the dotted red line). Following this, each object is assigned into a cluster. We applied t-SNE to project the high-dimensional data into a 2D space, using the cluster labels to colour the two-dimensional objects in the projected feature space (see Figure 2). We remark that there is fairly good agreement (following visual inspection) on the assigned clusters and the t-SNE 2D projection.


Figure 1: Dendrogram for the hierarchical clustering with Ward's linkage to determine the number of clusters in the analysis using all data. Following visual inspection we decided to opt for six clusters (highlighted with the dotted red line).


Figure 2: Two-dimensional representation of the original high-dimensional dataset using t-SNE and marking the six clusters (denoted C1…C6) computed using hierarchical clustering with the original feature set (see dendrogram in Figure 1).

i-Detect output

Figure 3: Output of the *i-Detect* algorithm assigning feature weights resulting in unsupervised feature selection.

Figure 3 presents the output of *i-Detect*, denoting the indices of the selected features associated with non-zero weights (the vast majority of the features were assigned zero weights and hence can be eliminated from further processing). We set a cut-off threshold at 0.05, which yielded 21 features. Overall, the selected feature subset comprises primarily wavelet-based features. We then repeated the process with hierarchical clustering (Figure 4) and 2D projection of the feature space spanned by the selected feature subset (Figure 5). We note that this time we decided on four clusters in the reduced feature space following visual inspection of the dendrogram, and again the 2D projection in Figure 5 is well aligned with the identified clusters. The computed four clusters were relatively evenly distributed with 458, 540, 577, and 522 objects in each.

Dendrogram, selected features

Figure 4: Dendrogram for the hierarchical clustering with Ward's linkage to determine the number of clusters in the analysis using all data. Following visual inspection we decided to opt for six clusters (highlighted with the dotted red line).

## t-SNE projection (selected features)

Figure 5: Two-dimensional representation of the dataset with selected features (seen in Figure 3) using t-SNE and marking the four clusters (denoted C1…C4) computed using hierarchical clustering with the selected feature subset (see Figure 4).

5 DISCUSSION

We explored the potential of processing features extracted using acoustic analysis of sustained vowel /a/ phonations in order to apply cluster analysis and define PD subtypes. Using unsupervised feature selection we determined a subset of 21 features from the originally high-dimensional subset of 307 features. We reported that the 2097 PD phonations used in the study could be clustered into four groups. Therefore, in principle a new PD participant could be phenotyped on the basis of a sustained vowel /a/ phonation to identify the PD group with which they are similar. In turn, if we could interpret what these clusters mean this may have important implications regarding PD symptom trajectory and developing better-targeted therapeutic strategies.

Interestingly, previous studies on PD subtyping have also reported the identification of four groups even though they had used very different data modalities. For example, Lewis et al. (2005), collected demographic, motor, mood, and cognitive measures from 120 early-stage PD participants and applied standard k-means cluster analysis which resulted into four main subgroups: (1) younger PD onset; (2) tremor-dominant; (3) non-tremor dominant with considerable cognitive impairment and mild depression; and (4) rapid disease progression but no cognitive impairment. Lawton et al. (2018) investigated motor, non-motor, and cognitive domains expressed using standardized clinical instruments on two large PD cohorts (1601 and 944 participants). They applied standard k-means

clustering on the latent variables extracted through factor analysis of the aggregate standardized questionnaires, and reported four main subgroups: (1) fast motor progression with symmetrical motor disease, poor olfaction, cognition and postural hypotension; (2) mild motor and non-motor disease with intermediate motor progression; (3) severe motor disease, poor psychological well-being and poor sleep with an intermediate motor progression; (4) slow motor progression with tremor-dominant, unilateral disease. van Rooden et al. (2011) similarly applied cluster analysis on two PD cohorts (344 and 357 participants) and reported four subgroups: (1) mildly affected in all domains, (2) predominantly severe motor complications, (3) affected mainly on nondopaminergic domains with no major motor complications, (4) severely affected across all domains. Mu et al. (2017) employed k-means domain clustering based on motor and non-motor symtoms in PD using two cohorts (411 and 540 participants), and similarly also reported finding four clusters: (1) mild, (2) non-motor dominant, (3) motor-dominant, and (4) severe.

Although there appear quite clear differences in the distributions of the selected features corresponding to each of the four clusters (results not shown) it is difficult to associate those with specific vocal performance degradation symptoms. In all cases, the wavelet coefficients used here correspond to expressing uncertainties in the F0. Moreover, it is not clear whether and how well the four identified clusters on the basis of the acoustic features extracted from the sustained vowel /a/ phonations match with the PD symptoms using in previous studies (Lewis et al. 2005; van Rooden et al., 2011; Lawton et al., 2018). Unfortunately, additional modalities or UPDRS assessments are not available in the PVI dataset, and other studies which have longitudinal clinical evaluations and patient reported outcome measures do not have speech signal recordings which would enable to explore bridging this gap.

The 2D projected feature space using t-SNE was intuitively appealing both when using the original high-dimensional dataset and also with the selected feature subset comprising 21 features: the clusters identified using hierarchical clustering appear to be generally well separated in the t-SNE derived scatter plots. This suggests that there is indeed some inherent underlying structure in the data, and that indeed the unsupervised feature selection algorithm has provided a feature subset that leads to some meaningful natural grouping of the PD cohort.

The field of PD subtyping on the basis of voice appears to have been scarcely investigated. Rueda

and Krishnan (2018) attempted cluster analysis algorithms on the basis sustained vowel /a/ recordings in 57 HC and 57 matched PD participants. However, the limited sample size suggests there is no sufficient statistical power to detect multiple clusters and hence their findings should be interpreted very tentatively. Moreover, mixing healthy controls with PD participants by design is not aimed to deliver PD subtypes but rather a more generic grouping of voices. We only used data from the PVI US cohort in this study. We decided to focus only on a single cohort to avoid potential language confounds in the design of cluster analysis; we are currently working on generalizing findings to the other cohorts in PVI, developing new insights when comparing derived cluster groups across the different locations where PD participants self-enrolled.

We envisage the PVI study and the findings presented herein may contribute towards improving understanding of the nature of PD subtypes and hence potentially informing therapeutic interventions in clinical practice (Triantafyllidis and Tsanas, 2019). We are further exploring the PVI data to investigate differences across PD cohorts at scale between different geographical locations, both towards understanding differences versus HC and also internal variability which may inform future clinical trials.

ACKNOWLEDGEMENTS

REFERENCES

Arora, S. Visanji, N.P., Mestre, T.A., Tsanas, A., Al Dakheel, A., Connolly, B.S., Gasca-Salas, C., Kern, D.S., Jain, J., Slow, E.J., Faust-Socher, A., Lang, A.E. Little, M.A., Marras C. 2018. Investigating voice as a biomarker for leucine-rich repeat kinase 2-associated Parkinson's disease: a pilot study, *Journal of Parkinson's Disease*, Vol. 8(4), pp. 503-510

Arora, S., Baghai-Ravary, L., Tsanas A. 2019. Developing a large scale population screening tool for the assessment of Parkinson's disease using telephone-quality speech, *Journal of Acoustical Society of America*, Vol. 145(5), 2871-2884

Bishop, C.M. 2006. *Pattern recognition and machine learning*, Springer

Camacho, A., Harris, J.G. 2008. A sawtooth waveform inspired pitch estimator for speech and music, *Journal of the Acoustical Society of America*, Vol. 124, 1638-1652

Dorsey, E.R., George, B.P., Leff, B., Willis A.W. 2013. The coming crisis: obtaining care for the growing burden of neurodegenerative conditions, *Neurology*, Vol. 80, 1989-1996

Duda, R.O., Hart, P.E., Stork, D.G. 2001. Pattern classification, Wiley-interscience, 2nd ed.

Dy, J.G., Brodley, C.E. 2004. Feature selection for unsupervised learning, *Journal of Machine Learning Research*, Vol. 5, 845-889

Fereshtehnejad, S.-M. et al. 2015. New clinical subtypes of parkinson disease and their longitudinal progression: a prospective cohort comparison with other phenotypes. *JAMA Neurology*, Vol. 72, pp. 863–873

GBD 2016 Parkinson's Disease Collaborators 2018. Global, regional, and national burden of Parkinson's disease, 1990–2016: a systematic analysis for the Global Burden of Disease Study 2016. *The Lancet Neurology*, Vol. 17, pp. 939-953

Gomez-Vilda, P., Mykyska, J., Gomez, A., Palacios, D., Rodellar, V., Alvarez A. 2019. Characterization of Parkinson's disease dysarthria in terms of speech articulation kinematics, *Biomedical Signal Processing and Control*, Vol. 52, 312-320

Guyon, I., Gunn, S., Nikravesh, M., Zadeh, L.A. (Eds.) 2006. *Feature Extraction: Foundations and Applications*, Springer

Hartelius L., Svensson P. 1994. Speech and swallowing symptoms associated with parkinson's disease and multiple sclerosis: A survey, *Folia Phoniatr. Logop.*, Vol. 46, pp. 9- 17

Hastie, T. Tibshirani, R. Friedman J. 2009. *The elements of statistical learning: data mining, inference, and prediction*, Springer, 2nd ed.

Horne, E., Tibble, H. Sheikh, A., Tsanas A. 2020. Challenges of clustering multimodal clinical data: a review of applications in asthma subtyping, *Journal of Medical Internet Research* (under review)

Lawton, M., Ben-Shlomo, Y., May, M.T., et al. 2018. Developing and validating Parkinson's disease subtypes and their motor and cognitive progression, *Journal of Neurology, Neurosurgery and Psychiatry*, Vol. 89, pp. 1279-1287

Lewis, S.J.G., Foltynie, T., Blackwell, A.D., Robbins, T.W., Owen, A.m., Barker R.A. 2005. Heterogeneity of Parkinson's disease in the early clinical stages using a data driven approach, *Journal of Neurology, Neurosurgery and Psychiatry*, Vol. 76, 343-348

Mu, J., Chaudhuri, K.R., Bielza, C., de Pedro-Cuesta, J., Larranaga, P., Martinez-Martin, P. 2017. Parkinson's disease subtypes identified from cluster analysis of motor and non-motor symptoms, *Frontiers in Aging Neuroscience*, 9:301

Nguyen, A., Roth, N., Ghassemi, N.H., Hannink, J., Seel, T., Klucken, J., Gassner, H., Eskofier, B.M. 2019. Development and clinical validation of inertial sensor-based gait-clustering methods in Parkinson's disease, *Journal of NeuroEngineering and Rehabilitation*, 16:77

Olanow, C.W., Stern, M.B., Sethi,K. 2009. The scientific and clinical basis forthe treatment of Parkinson disease, *Neurology*, Vol. 72 (21 Suppl 4) s1-s136

Rueda, A., Krishnan, S., 2018. Clustering Parkinson's and age-related voice impairment signal features for unsupervised learning, *Advances in Data Science and Adaptive Analysis*, Vol. 10(2);1840007

San Segundo, E., Tsanas, A., Gomez-Vilda, P., 2017. Euclidean distances as measures of speaker similarity including identical twin pairs: a forensic investigation using source and filter voice characteristics, *Forensic Science International*, Vol. 270, pp.25-38

Selikhova, M., Williams, D.R., Kempster, P.A., Holton, J.L., Revesz, T., Lees, A.J. 2009. A clinic-pathological study of subtypes in Parkinson's disease, *Brain*, Vol. 132, pp. 2947-2957

Sheaves, B., Porcheret, K., Tsanas, A., Espie, C., Foster, R., Freeman, D., Harrison, P.J., Wulff, K., Goodwin, G.M. 2016. Insomnia, nightmares, and chronotype as markers of risk for severe mental illness: results from a student population, *Sleep*, Vol. 39(1), pp. 173-181

Titze, I.R. 2000. *Principles of Voice Production*. National Center for Voice and Speech, Iowa City, US, 2nd printing

Triantafyllidis, A.K., Tsanas A. 2019. Applications of machine learning in real-life digital health interventions: review of the literature, *Journal of Medical Internet Research* (JMIR), Vol. 21(4), e12286

Tsanas, A., Little, M.A., McSharry, P.E., Ramig, L.O. 2010. New nonlinear markers and insights into speech signal degradation for effective tracking of Parkinson's disease symptom severity, *International Symposium on Nonlinear Theory and its Applications* (NOLTA), pp. 457-460, Krakow, Poland, 5-8 September

Tsanas, A., Little, M.A., McSharry, P.E., Ramig, L.O. 2011. Nonlinear speech analysis algorithms mapped to a standard metric achieve clinically useful quantification of average Parkinson's disease symptom severity, *Journal of the Royal Society Interface*, Vol. 8, pp. 842-855

Tsanas, A., Little, M.A., McSharry, P.E., Spielman, J., Ramig, L.O. 2012. Novel speech signal processing algorithms for high-accuracy classification of Parkinson's disease, *IEEE Transactions on Biomedical Engineering*, Vol. 59, 1264-1271

Tsanas A. 2012. *Accurate telemonitoring of Parkinson's disease symptom severity using nonlinear speech signal processing and statistical machine learning*, Ph.D. thesis, Oxford Centre for Industrial and Applied Mathematics, University of Oxford

Tsanas A. 2013. Acoustic analysis toolkit for biomedical speech signal processing: concepts and algorithms, 8th International *Workshop on Models and Analysis of Vocal Emissions for Biomedical Applications* (MAVEBA), pp. 37-40, Florence, Italy, 16-18 December

Tsanas, A., Gómez-Vilda P., 2013. Novel robust decision support tool assisting early diagnosis of pathological voices using acoustic analysis of sustained vowels, *Multidisciplinary Conference of Users of Voice, Speech and Singing* (JVHC 13), pp. 3-12, Las Palmas de Gran Canaria, 27-28 June

Tsanas, A. Little, M.A., McSharry P.E. 2013. A methodology for the analysis of medical data, in *Handbook of Systems and Complexity in Health*, Eds. J.P. Sturmberg, and C.M. Martin, Springer, pp. 113-125 (chapter 7)

Tsanas, A., Little, M.A., Fox, C., Ramig L.O. 2014a. Objective automatic assessment of rehabilitative speech treatment in Parkinson's disease, *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, Vol. 22, 181-190

Tsanas, A., Zañartu, M., Little, M.A., Fox, C., Ramig, L.O., Clifford, G.D. 2014b. Robust fundamental frequency estimation in sustained vowels: detailed algorithmic comparisons and information fusion with adaptive Kalman filtering, *Journal of the Acoustical Society of America*, Vol. 135, 2885-2901

Tsanas A. 2019. New insights into Parkinson's disease through statistical analysis of standard clinical scales quantifying symptom severity, *41st IEEE Engineering in Medicine in Biology* Conference (EMBC), Berlin, Germany, 23-27 July

Tsanas, A., Arora S. 2019. Biomedical speech signal insights from a large scale cohort across seven countries: the Parkinson's voice initiative study, *Models and Analysis of Vocal Emissions for Biomedical Applications*, Florence, Italy, 17-19 December

Yao, J., Mao, Q., Goodison, S., Mai, V., Sun, Y. 2015. Feature selection for unsupervised learning through local learning. *Pattern Recognition Letters*, Vol. 53, pp. 100–107

van der Maaten, L.J.P., Hinton, G.E. 2008. Visualizing high-dimensional data using t-SNE, *Journal of Machine Learning Research*, Vol. 9, pp. 2579-2605

van Rooden S.M., Colas F., Martínez-Martín P., Visser M., Verbaan D., Marinus J., Chaudhuri R.K., Kok J.N., van Hilten J.J. 2011. *Movement Disorders*, Vol. 26(1), pp.51-58.

Zhang, X., Chou, J., Liang, J., Xiao, C., Zhao, Y., Sarva, H., Henchcliffe, C., Wang, F. 2019. Data-driven subtyping of Parkinson's disease using longitudinal clinical records: a cohort study, *Scientific Reports*, Vol. 9;797.